

Real-Time American Sign Language Detection Using MediaPipe Hand Localization and ResNet-50 Classification

Udit Asthana

Manipal Institute of Technology

Reg. No: 240905310

Roll No: 31

Abhishek L Prabhu

Manipal Institute of Technology

Reg. No: 240905676

Roll No: 71

Ashwath Patil

Manipal Institute of Technology

Reg. No: 240905460

Roll No: 46

Kishan Kanvathirtha

Manipal Institute of Technology

Reg. No: 240905678

Roll No: 72

Arnav Dugad

Manipal Institute of Technology

Reg. No: 230905227

Roll No: 73

Abstract—American Sign Language (ASL) is the primary communication mode for the hearing-impaired community. Real-time automated recognition of ASL gestures remains an open engineering problem. This paper describes a two-stage pipeline: hand region localization via MediaPipe, followed by letter classification with a fine-tuned ResNet-50 backbone and a custom head. Training uses the publicly available ASL Alphabet dataset from Kaggle, evaluated over 26 letter classes. A cosine annealing scheduler with linear warm-up, gradient clipping, and early stopping are applied during training. Test accuracy: 99.80%. The system runs in real-time on a live webcam feed. Hand skeleton overlays and predicted letter labels are rendered per frame. A FastAPI inference server and browser frontend enable interactive deployment. Landmark-based detection combined with CNN classification is viable for low-latency sign language recognition.

Index Terms—sign language recognition, ASL, MediaPipe, ResNet-50, hand detection, real-time inference, deep learning, convolutional neural networks

I. Introduction

Sign language is the primary communication mode for over 70 million deaf and hard-of-hearing people worldwide [1]. ASL specifically — used mostly in North America — uses hand gestures and finger-spelling to express words and concepts. Despite how widely it’s used, there are almost no real-time automated translation tools available. That gap makes everyday communication with the hearing majority significantly harder than it needs to be.

CNNs and real-time object detection frameworks have opened up new options for gesture recognition. Systems that can classify a gesture from a live video frame in milliseconds are a meaningful step toward accessible communication technology. Our approach uses two stages: first, a MediaPipe hand tracking module localizes the hand in each camera frame, producing a tight bounding box crop (not visible to the user) and a skeleton overlay (visible). Second, a fine-tuned ResNet-50 [3] classifies the cropped

image into one of 26 ASL alphabet letters. The result is overlaid on the live stream in real time.

Our contributions:

- A modular training pipeline with configurable optimizer, scheduler, and early stopping, implemented in PyTorch.
- MediaPipe-based hand detection integrated with a ResNet-50 classification head for end-to-end ASL letter recognition.
- A real-time inference module with webcam integration and visual feedback.
- A FastAPI web interface for browser-based interaction with the live pipeline.

II. Related Work

Early sign language recognition systems depended on sensor-based hardware: data gloves, depth cameras, 3D pose capture [4]. These methods worked. They also imposed hardware constraints that blocked practical deployment.

Deep learning shifted the dominant approach to image-based methods. Convolutional networks trained on static hand images showed strong performance on isolated gesture classification [5]. Transfer learning from ImageNet-pretrained backbones — VGG-16, ResNet, MobileNet — extended this to smaller, domain-specific datasets with improved accuracy.

Real-time landmark-based hand tracking has been applied to gesture spotting. End-to-end two-stage systems covering the full ASL alphabet in a live-camera setting are less common. This work targets that gap.

III. Dataset

The ASL Alphabet dataset [6] was found on Kaggle. It contains 87,000 labeled images across 29 classes: 26 letters (A–Z), plus SPACE, DELETE, and NOTHING.

Images are resized to 224×224 to match the ResNet-50 pre-trained weights. Normalization was done using standard ImageNet scheme and settings. The dataset is partitioned randomly: 70% training, 20% validation, 10% test using random-split. Random splitting preserves class balance across split.

A `_TransformSubset` wrapper applies separate transforms to each split. To increase data diversity, Training images were flipped randomly, given jitters etc.

Validation and test images receive only deterministic resize and normalization. Batch size was set 128 to accommodate for resource and compute constraints. Shuffling is enabled for training and disabled for evaluation.

Approximate split sizes: 60,900 training samples, 17,400 validation samples, 8,700 test samples.

IV. Methodology

A. System Overview

Inference proceeds in two sequential stages.

- 1) Hand Detection (MediaPipe): Each video frame is passed to MediaPipe. Bounding boxes are returned around detected hand regions.
- 2) Letter Classification (ResNet-50): The hand crop is extracted, resized, and forwarded to the classifier. A probability distribution over 26 ASL letter classes is returned.

The predicted letter and bounding box are overlaid on the original frame.

B. Hand Detection Module

The `HandDetector` class wraps the MediaPipe `HandLandmarker` Tasks API. The operating mode is `VIDEO`. This mode enables frame-to-frame tracking through MediaPipe’s internal Kalman filter.

The Kalman filter requires strictly monotonically increasing real wall-clock millisecond timestamps. Passing bare integer counters (1, 2, 3, ...) is a common implementation error. It mis-models velocity. Landmark jitter increases. Timestamps here are derived from `time.monotonic_ns()`.

Detection thresholds are raised above MediaPipe defaults. Minimum hand detection and presence confidence: 0.65 (default: 0.50). Tracking confidence: 0.55. The defaults are too permissive for static-sign ASL. Partially-occluded hands during letter transitions produce false detections under the defaults.

Per-frame output is a list of two hand results. For each hand, the raw landmark bounding box is computed from the min/max of all 21 normalized landmark coordinates. The box is expanded by 35% per dimension (expansion=0.35). This gives environmental context to the images of hands detected. This makes the job of classification easier for the Pre-Trained Resnet Backbone. The expanded box is then smoothened across frames via exponential smoothing: $0.15 \cdot \text{prev} + 0.85 \cdot \text{new}$.

Motion magnitude is computed on the raw (pre-smoothing) center against the previous smoothed center. Motion is measured as Euclidean (L2) distance. Manhattan (L1) distance under-reports diagonal motion by approximately 30%.

Per-hand tracking is by keying hand-index. A single shared variable fails silently when two hands appear in the same frame. When a hand goes out of frame, its tracking state vanishes with it. The next detection starts from scratch. Zero-area ROI crops after boundary clamping are thrown away; they are not passed to the classifier. When multiple hands are detected at inference time, the detection with the highest classifier score is selected:

$$\text{det}^* = \arg \max_{d \in \mathcal{D}} d.\text{score} \quad (1)$$

C. Classification Model (SLD)

The Sign Language Detector (SLD) uses a ResNet-50 backbone. All backbone parameters were frozen, meaning gradient tracking for them was disabled. A custom trainable classification head was added to the backbone.

The backbone’s final average pooling layer outputs a 2048-dimensional feature vector. The original fully-connected layer is discarded entirely. The replacement head is:

$$\hat{y} = \text{head}(\text{backbone}(x)) \quad (2)$$

$$\text{head}(z) = W_2 \cdot \text{GELU}(\text{Dropout}(\text{LayerNorm}(z) \cdot W_1)) \quad (3)$$

$W_1 \in \mathbb{R}^{2048 \times 512}$, $W_2 \in \mathbb{R}^{512 \times 29}$. Output covers 29 classes: 26 letters plus del, nothing, space.

LayerNorm is used at the backbone-to-head boundary. BatchNorm is not. BatchNorm is sensitive to batch size. Its behavior differs between train and eval modes. With a frozen backbone and a trainable head only, this instability is unacceptable. LayerNorm normalizes per-sample. It is invariant to batch size.

GELU is used in place of ReLU. GELU is smooth. Its gradient is non-zero at negative values. Dropout at $p = 0.3$ regularizes the intermediate 512-dimensional representation.

At inference, the checkpoint is loaded via `torch.load` with `weights_only=False`.

This is explicit. The checkpoint stores optimizer and scheduler state as well as model weights. It cannot be loaded in `weights-only` mode.

D. Training Configuration

Training configuration is kept in a `dataclass` (`TRAINING_CONFIG`).

The dataset is split apart via `random_split`: 70% training, 20% validation, 10% test. A `_TransformSubset` wrapper is applied to each split. This allows training and evaluation transforms to differ without modifying

the underlying dataset object. The naive alternative – overwriting `subset.dataset` – causes the Subset’s indices to point into a wrongly-sized dataset, raising an `IndexError`. Training images receive augmentation. Validation and test images receive only deterministic resize and normalization.

Per-batch metrics – training loss, accuracy, gradient norm – are logged to W&B at each step via `wandb.log(..., step=self.global_step)`. At epoch level, mean gradient norm and gradient clipping ratio are additionally reported. The clipping ratio is the fraction of batches where the gradient norm exceeded the clip threshold.

TABLE I
Training Hyperparameters

Parameter	Value
Optimizer	SGD (Nesterov)
Learning Rate	1×10^{-2}
Momentum	0.9
Weight Decay	5×10^{-4}
Scheduler	Cosine Annealing + Warm-Up
Warm-Up Epochs	10
$T_{\max}$	matched to epoch count
$\eta_{\min}$	1×10^{-4}
Max Epochs	200 (3 in practice)
Batch Size	128
Gradient Clip Norm	1.0
Early Stopping Patience	20

Training ran for 3 epochs. Google Colab session limits constrained this. The configured maximum was 200. `T_max` and epoch count were overridden via command-line at runtime. Results are discussed in Section VI.

1) Learning Rate Schedule: A cosine annealing schedule with linear warm-up is used. Warm-up spans epochs 0 to $T_w - 1$. During warm-up, the learning rate increases linearly:

$$\eta_t = \eta_{\text{base}} \cdot \frac{t + 1}{T_w} \quad (4)$$

After warm-up, cosine decay applies:

$$\eta_t = \eta_{\min} + \frac{\eta_{\text{base}} - \eta_{\min}}{2} \left(1 + \cos \left(\pi \cdot \frac{t - T_w}{T_{\max} - T_w} \right) \right) \quad (5)$$

Early instability is suppressed by the warm-up phase. The learning rate settles toward $\eta_{\min}$ near convergence.

2) Gradient Clipping: Gradient norm clipping is applied after each backward pass. Threshold: $\tau = 1.0$.

$$g \leftarrow g \cdot \min \left(1, \frac{\tau}{\|g\|_2} \right) \quad (6)$$

3) Early Stopping: This was an implemented feature which did not come in handy, due to compute limitations.

E. Experiment Tracking

Weights & Biases (W&B) [7] was used to log training and validation metrics. Per-batch metrics include training loss, accuracy, gradient norm. validation loss and accuracy,

learning rate, mean gradient norm were logged per epoch, along with gradient clipping ratio. Test loss and accuracy are logged as summary metrics at run end.

V. Real-Time Inference Pipeline

At the time of scanning, frames are captured online via a live webcam feed and passed into a classification pipeline that follows this structure

- 1) Hand Detection: Media pipe detects the hand and contextual environment.
- 2) Stability Check: Euclidean motion between consecutive frames is used to flag instability
- 3) Crop and Pad: This is invisible to the user, the box that establishes context is not visible. However, the skeleton tracing the landmarks of the detected hand is available to the user.
- 4) Classification: The crop is resized to 224×224 and normalized. It is passed through the SLD model which predicts several possible classifications. A prediction is accepted only if confidence exceeds 0.50.
- 5) Temporal Voting: A majority-vote buffer of size 10 is maintained. A letter is emitted only when 70% or more of the buffer agrees.
- 6) Cooldown: A 1-second cooldown is enforced. The same letter cannot be emitted repeatedly in rapid succession.

A. FastAPI Inference Server

The Inference Pipeline is available through a lightweight FastAPI [8] server as a REST endpoint. Any HTTP client can send JPEG images to the server, while JSON output will be received.

Endpoint: `POST /predict` accepts a multipart-encoded JPEG image. The response carries:

- letter — stable, voted prediction (A–Z, del, nothing, space). null if not yet stable.
- confidence — softmax probability of the top class.
- bbox — bounding box as $[x, y, w, h]$ in pixel coordinates.
- handedness — left or right, from MediaPipe.
- motion — Euclidean motion magnitude in pixels. Used for stability gating.
- landmarks — 21 normalized hand landmark coordinates. Used for frontend skeleton rendering.
- error — status string (no hand, hand moving, low confidence, not stable yet, cooldown) when no letter is emitted.

The Global State consists of several locks, namely `threading.Lock`, including those related to the Prediction Buffer, Last Emitted Letter, and Cooldown Time Stamp. CORS is enabled for all origins.

B. Browser Frontend

Single-page HTML/JS frontend, connects to the FastAPI server. Client-side MediaPipe runs through the Tasks Vision JS SDK at 60 fps – for the skeleton overlay,



Fig. 1. FastAPI /predict endpoint response for a live hand frame, showing letter, confidence, bounding box, handedness, motion, and landmark fields.

which renders with zero latency. The actual classification frames go to the backend at around 8 fps.

For letter emission we use a dwell-based approach. A letter has to stay stable for 1.8s before it gets added to the transcript. If the same letter repeats, it needs 3.6s. This way brief hand holds don’t accidentally trigger output.

VI. Results and Discussion

Training was conducted on a Google Colab T4 GPU. Session time limits constrained fine-tuning to 3 epochs. The configured maximum was 200. Results below reflect this constrained run.

The model was trained for 3 epochs. Batch size: 128. Optimizer: SGD with Nesterov momentum. Scheduler: cosine annealing with linear warm-up, as described in Section IV. Final performance metrics are summarized in Table II.

TABLE II
Training and Evaluation Results

Split	Loss	Accuracy (%)
Train (epoch 3)	0.0474	99.03
Validation	0.0099	99.86
Test	0.0096	99.80

Test accuracy reached 99.80% at epoch 3. Validation accuracy at epoch 1 was 64.41%. By epoch 3 it reached 99.86%. Training was terminated well short of the configured limit.

These numbers require qualification. The ImageNet-pretrained backbone carries a strong prior. The dataset

is constrained: controlled backgrounds, consistent framing, one hand per image. High accuracy under these conditions is not a claim of general robustness.

W&B logs at the final epoch show a mean gradient norm of 1.64 and a clipping ratio of 0.77. The clipping constraint was still active and the loss had not yet flattened. Training to full convergence would likely yield marginal gains on this dataset; the benefit would be more meaningful on harder, out-of-distribution evaluation sets.

The main challenge for deployment is distribution shift. The training set does not cover varied skin tones, uncontrolled lighting, or cluttered backgrounds. The two-stage design helps partially – MediaPipe isolates the hand crop before it reaches the classifier, which cuts down on background interference.

A. Colab Cell Output

The Colab cell output is attached for reference. The full implementation is available at the official GitHub repository: github.com/madhavasthana-oss/SLD_MAHE.

VII. Conclusion

This paper presents a real-time ASL alphabet recognition system. The two-stage design — MediaPipe hand detection followed by ResNet-50 classification — achieves 99.80% test accuracy on the Kaggle ASL Alphabet dataset. The system operates at interactive frame rates. Deployment is via a FastAPI inference server and browser-based frontend.

The training pipeline uses warm-up scheduling, Nesterov SGD, gradient clipping, and early stopping. A dwell-




Fig. 2. Training and validation loss/accuracy curves across steps along with gradient norm, logged via Weights & Biases.

based temporal emission mechanism and majority-vote buffer stabilize letter output during live use. The system addresses a real accessibility problem. It is a functional baseline, not a ceiling.

Future directions:

- Fine-tuning MediaPipe on a dedicated hand detection dataset for improved localization.
- Upgrading the backbone to ResNet-152 or EfficientNet for higher classification capacity.
- Extending to dynamic gestures and full ASL words via temporal models (LSTM, Transformer).
- Mobile deployment for broader accessibility reach.

References

- [1] World Health Organization, “Deafness and hearing loss,” WHO Fact Sheets, 2021. [Online]. Available: <https://www.who.int/news-room/fact-sheets/detail/deafness-and-hearing-loss>
- [2] J. Redmon, S. Divvala, R. Girshick, and A. Farhadi, “You only look once: Unified, real-time object detection,” in Proc. IEEE CVPR, 2016, pp. 779–788.
- [3] K. He, X. Zhang, S. Ren, and J. Sun, “Deep residual learning for image recognition,” in Proc. IEEE CVPR, 2016, pp. 770–778.
- [4] S. Mitra and T. Acharya, “Gesture recognition: A survey,” IEEE Trans. Syst., Man, Cybern. C, Appl. Rev., vol. 37, no. 3, pp. 311–324, 2007.
- [5] L. Pigou, S. Dieleman, P.-J. Kindermans, and B. Schrauwen, “Sign language recognition using convolutional neural networks,” in Proc. ECCV Workshops, 2015.
- [6] A. Akash, “ASL Alphabet,” Kaggle, 2018. [Online]. Available: <https://www.kaggle.com/datasets/grassknoted/asl-alphabet>
- [7] L. Biewald, “Experiment tracking with weights and biases,” 2020. [Online]. Available: <https://www.wandb.com>
- [8] S. Ramírez, “FastAPI,” 2018. [Online]. Available: <https://fastapi.tiangolo.com>